

# 1.Introduction

## About Walmart

Walmart is an American multinational retail corporation that operates a chain of supercenters, discount departmental stores, and grocery stores from the United States. Walmart has more than 100 million customers worldwide.

## Business Problem

The Management team at Walmart Inc. wants to analyze the customer purchase behavior (specifically, purchase amount) against the customer's gender and the various other factors to help the business make better decisions. They want to understand if the spending habits differ between male and female customers: Do women spend more on Black Friday than men?

## 2.0 Exploratory Data Analysis

```
df.sample(10)
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
203754	1001447	P00116542	M	18-25	4	A	0
309709	1005738	P00182342	M	26-35	12	A	1
29345	1004458	P00010842	F	26-35	4	B	1
3660	1000594	P00192942	F	55+	13	B	3
95406	1002776	P00302742	M	26-35	14	C	4+
278003	1000881	P00186242	M	18-25	14	A	1
273223	1000081	P00114042	F	26-35	0	A	1
178599	1003618	P00043542	M	55+	17	A	4+
416932	1004124	P00226342	M	26-35	18	C	1
534187	1004250	P00189142	F	51-55	0	C	4+

**Observations** - - We got 10 columns and from the look of things, the only column we can consider right as a Numerical variable type is 'Product Column'. - All other columns are under the categorical variable category.

**Required Action** - - Checking if the columns have the correct datatype and if not, the respective columns needed to be changed to their appropriate datatype.

```
df.shape
```

```
(550068, 10)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                            550068 non-null  object
2   Gender                                550068 non-null  object
3   Age                                    550068 non-null  object
4   Occupation                             550068 non-null  int64
5   City_Category                         550068 non-null  object
6   Stay_In_Current_City_Years            550068 non-null  object
7   Marital_Status                        550068 non-null  int64
8   Product_Category                      550068 non-null  int64
9   Purchase                              550068 non-null  int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB
```

**Observation** - - The structure of the data frame consist of - 550068 columns and - 10 rows - The data types for the columns are not properly assigned.

**Further Action** - - Each column should be assigned the appropriate datatype to suit out analysis.

```
df_unique_check = df.copy()
for i in df.columns:
    if i == 'Purchase':
        continue
    df[i] = df[i].astype('category')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  category
1   Product_ID                            550068 non-null  category
2   Gender                                550068 non-null  category
3   Age                                    550068 non-null  category
4   Occupation                             550068 non-null  category
```

```

5   City_Category      550068 non-null  category
6   Stay_In_Current_City_Years  550068 non-null  category
7   Marital_Status     550068 non-null  category
8   Product_Category   550068 non-null  category
9   Purchase           550068 non-null  int64
dtypes: category(9), int64(1)
memory usage: 10.3 MB

```

**Observations** - - All the columns except for the Purchase column is changed to “Categorical Varibale”.

### **Further Action -**

*Handling duplicates and analysis of Unique values -*

```
df.nunique()
```

```

User_ID      5891
Product_ID   3631
Gender        2
Age           7
Occupation   21
City_Category 3
Stay_In_Current_City_Years  5
Marital_Status  2
Product_Category  20
Purchase     18105
dtype: int64

```

**Observations** - - From the above line what we can understand that despite having 550068 rows we have just 5891 User\_ID records. - This implies that this is a sales data where a single user will have multiple records which is based on their buying frequency. - Generally, with these kinds of data there are high chances of duplicate records since we don’t have the sales\_id to differentiate each sales.

**Further Action** - - Find and remove the duplicates. - Make a summary with unique values to understand our dataset.

*Finding the Unique value in each column*

```

def unique_values():
    for i in df_unique_check.columns:
        if df_unique_check[i].dtype == 'int64':
            print(f"Unique values in {i} column are")
            print(np.sort(df_unique_check[i].unique()), end='\n\n')
        else:
            print(f"Unique values in {i} column are")

```

```
print(df_unique_check[i].unique(), end='\n\n')
unique_values()
```

Unique values in User\_ID column are  
[1000001 1000002 1000003 ... 1006038 1006039 1006040]

Unique values in Product\_ID column are  
['P00069042' 'P00248942' 'P00087842' ... 'P00370293' 'P00371644'  
'P00370853']

Unique values in Gender column are  
['F' 'M']

Unique values in Age column are  
['0-17' '55+' '26-35' '46-50' '51-55' '36-45' '18-25']

Unique values in Occupation column are  
[ 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20]

Unique values in City\_Category column are  
['A' 'C' 'B']

Unique values in Stay\_In\_Current\_City\_Years column are  
['2' '4+' '3' '1' '0']

Unique values in Marital\_Status column are  
[0 1]

Unique values in Product\_Category column are  
[ 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20]

Unique values in Purchase column are  
[ 12 13 14 ... 23959 23960 23961]

**Observations** - - The unique values are consolidated and sorted based on their datatype for the ease of analysis.

```
df['Marital_Status'] = df['Marital_Status'].replace({0:'Single', 1:'Married'})
```

*Finding if the columns are properly structured and making sure that there are no nested values*

```
df[df.duplicated()].value_counts()
```

Series([], Name: count, dtype: int64)

**Observations** - - There are no duplicate records and the dataset is complete.

**Further Action** - - Checking the structure of the dataset.

*Checking for Nested Values*

```
def nested_values_check():
    df_check = df.copy() # Creating a copy to have the main dataframe unaffected.
    for i in df_check.columns:
        df_check[i] = df_check[i].astype('str')
        if i == 'Purchase':
            if df_check[i].str.contains(',').any()==False:
                print(f"{i} column--> Structured Properly")
            else:
                print(f"{i} column--> has Nested values")
        else:
            if df_check[i].str.contains(',').any()==False:
                print(f"{i} column--> Structured Properly", end = '\n\n')
            else:
                print(f"{i} column--> has Nested values", end = '\n\n')

nested_values_check()
```

User\_ID column--> Structured Properly

Product\_ID column--> Structured Properly

Gender column--> Structured Properly

Age column--> Structured Properly

Occupation column--> Structured Properly

City\_Category column--> Structured Properly

Stay\_In\_Current\_City\_Years column--> Structured Properly

Marital\_Status column--> Structured Properly

Product\_Category column--> Structured Properly

Purchase column--> Structured Properly

- It is clear that the data in our dataset is clean and structured properly to suit our analysis.

```
df.describe(include=['category']).T
```

	count	unique	top	freq
User_ID	550068	5891	1001680	1026
Product_ID	550068	3631	P00265242	1880
Gender	550068	2	M	414259
Age	550068	7	26-35	219587
Occupation	550068	21	4	72308
City_Category	550068	3	B	231173
Stay_In_Current_City_Years	550068	5	1	193821
Marital_Status	550068	2	Single	324731
Product_Category	550068	20	5	150933

```
df.describe().T
```

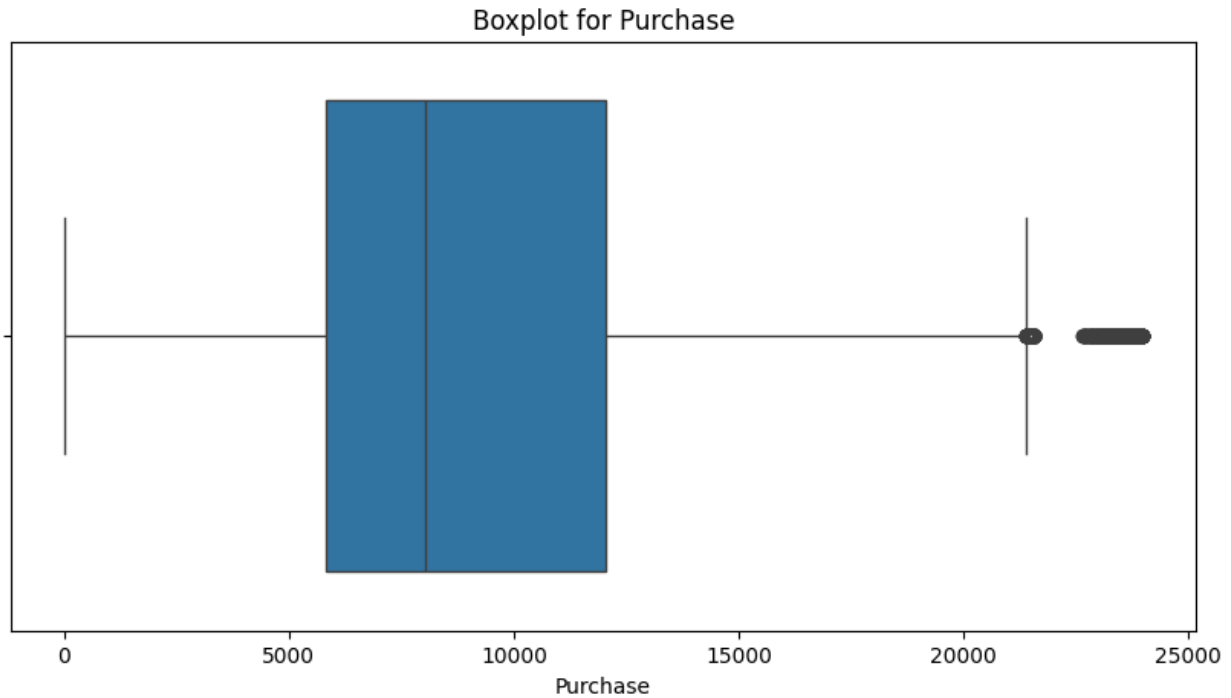
	count	mean	std	min	25%	50%	75%	max
Purchase	550068.0	9263.968713	5023.065394	12.0	5823.0	8047.0	12054.0	23961.0

## 3.Univariate Analysis

### 3.1 Numerical Variables

#### 3.1.1 Detecting Outliers

```
plt.figure(figsize=(10,5))
sns.boxplot(data = df, x = 'Purchase')
plt.title('Boxplot for Purchase')
plt.show()
```



*Calculating the number of Outliers*

```
# Calculating the Quartiles
Q1 = df['Purchase'].quantile(0.25)
Q3 = df['Purchase'].quantile(0.75)

# Calculating the IQR
IQR = Q3 - Q1

# Setting lower and upper limits to sepperrate the Outliers
lower_limit = Q1 - 1.5 * IQR
upper_limit = Q3 + 1.5 * IQR

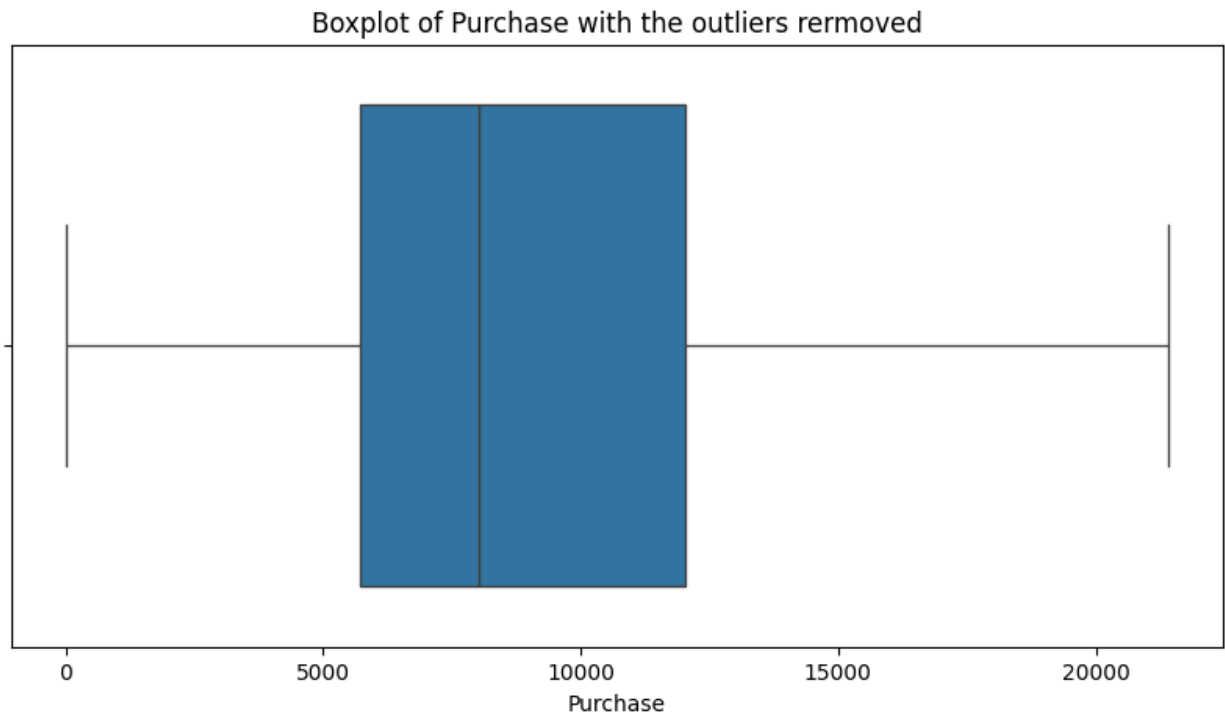
# Calculating the total outliers
lower_outliers = df[df['Purchase'] < lower_limit]
upper_outliers = df[df['Purchase'] > upper_limit]
Total_Outliers = len(lower_outliers) + len(upper_outliers)
print(f'Total number of outliers in our dataset is {Total_Outliers}')
```

Total number of outliers in our dataset is 2677

*Removing the Outliers*

```
remove_outliers = df[(df['Purchase'] > lower_limit) & (df['Purchase'] <
upper_limit)]
```

```
plt.figure(figsize=(10,5))
sns.boxplot(data = remove_outliers, x = 'Purchase')
plt.title('Boxplot of Purchase with the outliers removed')
plt.show()
```

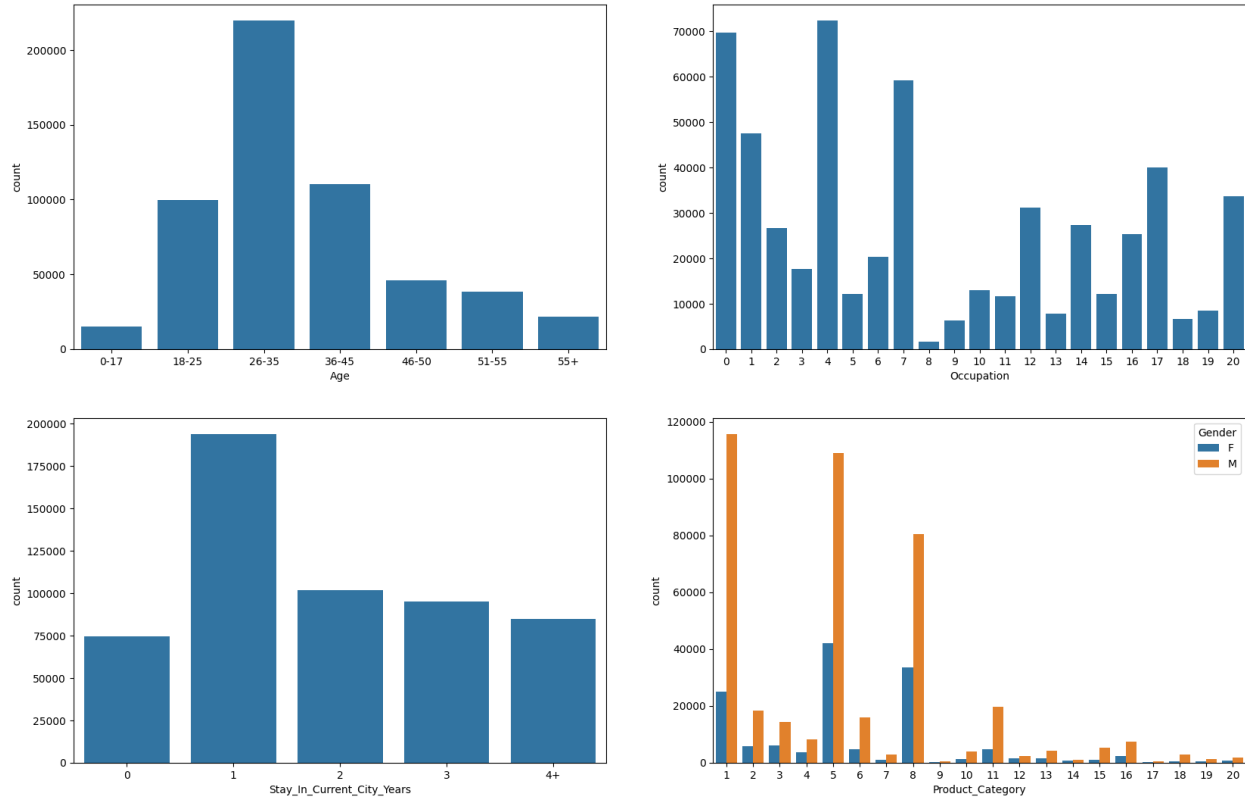


### 3.2 Categorical Variables

```
fig, axs= plt.subplots(2,2 , figsize=(20,13))
sns.countplot(data=df,x='Age',ax=axs[0,0])
sns.countplot(data=df,x='Occupation',ax=axs[0,1])
sns.countplot(data=df,x='Stay_In_Current_City_Years',ax=axs[1,0])
sns.countplot(data=df,x='Product_Category',ax=axs[1,1], hue='Gender')

plt.show()
```





*Insights* - Age between 26-35 were quite active and spent a significant amount during black friday sale. - People from the occupation 0, 4 and 7 must have quite a spending power as they were the top buyers of that category. - In the product category people preferred 1,5 and 8 more than any other product and the gender don't bear much impact since both gender have almost the same preference during the sale. - The people who are staying in the city for a 1 year took good advantage of the black friday sale.

```
# Finding Proportion of City, Gender and Marrital_status in the given data
City_prop = pd.DataFrame(df.groupby('City_Category')['User_ID'].count())
C_values = City_prop['User_ID'].tolist()
C_labels = City_prop.index

Gender_prop = pd.DataFrame(df.groupby('Gender')['User_ID'].count())
G_values = Gender_prop['User_ID'].tolist()
G_labels = Gender_prop.index

Marital_prop = pd.DataFrame(df.groupby('Marital_Status')['User_ID'].count())
M_values = Marital_prop['User_ID'].tolist()
M_labels = Marital_prop.index


plt.subplot(1,3,1)
plt.pie(G_values, labels = G_labels, autopct = '%0.2f%%')

plt.subplot(1,3,2)
```

```
plt.pie(M_values, labels = M_labels, autopct = '%0.2f%%')

plt.subplot(1,3,3)
plt.pie(C_values, labels = C_labels, autopct = '%0.2f%%')

plt.show()
```



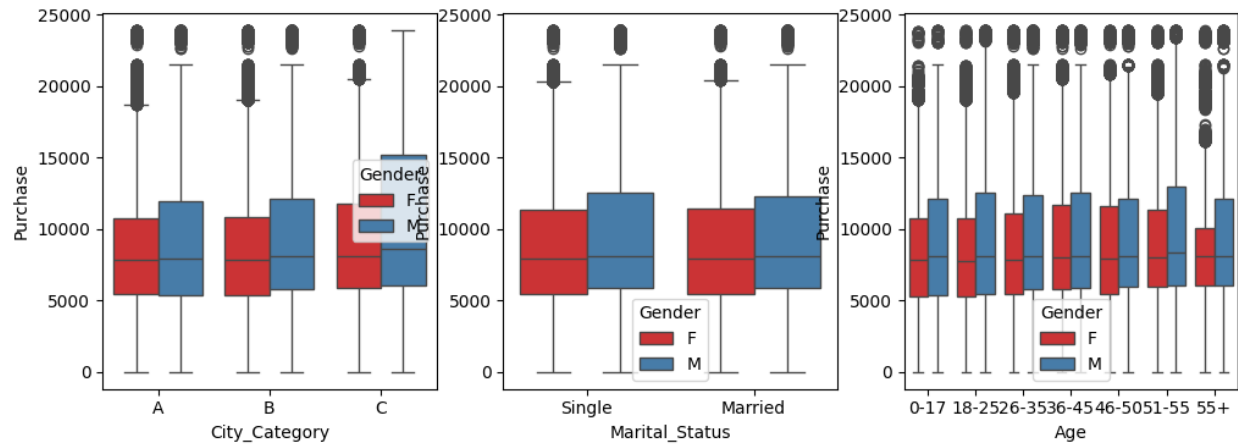
*Insights* - Most of the buyers are Men and single. - People from the Area B are quite active during this Black Friday.

## 4. Multivariate Analysis

### 4.1 Purchase vs different aspects of customers

```
fig, axis = plt.subplots(1,3, figsize=(12,4))
sns.boxplot(df, x="City_Category", y="Purchase", ax=axis[0], hue= 'Gender',
palette='Set1')
sns.boxplot(df, x="Marital_Status", y="Purchase", ax=axis[1], hue= 'Gender',
palette='Set1')
sns.boxplot(df, x="Age", y="Purchase", ax=axis[2], hue= 'Gender', palette='Set1')
plt.show()


# sns.boxplot(df, x="City_Category", y="Purchase", hue= 'Gender', palette='Set1')
# plt.show()
```



*Insights* - From the above chart we can see that the spending habit of both Genders regardless of their age, Marital status or the city they are from, all have relatively a same spending pattern and maintains the spending average around 8000\$.


## 4.2 Different Age groups and their buying habits

```
sns.histplot(df, x="Product_Category", hue="Age", multiple="stack",
palette='Set1')
plt.show()
```



### 4.3 Gender and thier Product preference

```
sns.histplot(df, x="Product_Category", hue="Gender", multiple="stack",  
palette='Set1')  
plt.show()
```



## 5. Confidence Interval

```
def df_ci(column, confidence_level):
    male_data = df[df[column] == 'M']['Purchase']
    male_mean = male_data.mean()
    male_std = male_data.std()
    male_se = male_std/np.sqrt(len(male_data))

    male_z = norm.ppf((1 + confidence_level)/2)
    male_ci = [male_mean - (male_z * male_se), male_mean + (male_z * male_se)]
    male_range = male_ci[1] - male_ci[0]

    female_data = df[df[column] == 'F']['Purchase']
    female_mean = female_data.mean()
    female_std = female_data.std()
    female_se = female_std/np.sqrt(len(female_data))

    female_z = norm.ppf((1 + confidence_level)/2)
    female_ci = [female_mean - (female_z * female_se), female_mean + (female_z *
female_se)]
    female_range = female_ci[1] - female_ci[0]
```

```

print(f'Records of {column} column')
print(f'Male {confidence_level*100}% confidence interval:{male_ci}')
print(f'Male {confidence_level*100}% range:{male_range}')
print(f'Female {confidence_level*100}% confidence interval:{female_ci}')
print(f'Female {confidence_level*100}% range:{female_range}')
print('\n')

```

```

def sample_ci(column, confidence_level, sample_size):
    male_data = df[df[column] == 'M']['Purchase']
    sample_male = male_data.sample(sample_size, random_state = 42)
    male_mean = sample_male.mean()
    male_std = sample_male.std()
    male_se = male_std/np.sqrt(len(sample_male))

    male_z = norm.ppf((1 + confidence_level)/2)
    male_ci = [male_mean - (male_z * male_se), male_mean + (male_z * male_se)]
    male_range = male_ci[1] - male_ci[0]

    female_data = df[df[column] == 'F']['Purchase']
    sample_female = female_data.sample(sample_size, random_state = 42)
    female_mean = sample_female.mean()
    female_std = sample_female.std()
    female_se = female_std/np.sqrt(len(sample_female))

    female_z = norm.ppf((1 + confidence_level)/2)
    female_ci = [female_mean - (female_z * female_se), female_mean + (female_z *
female_se)]
    female_range = female_ci[1] - female_ci[0]

    print(f'{sample_size} samples of {column} column')
    print(f'Male {confidence_level*100}% confidence interval:{male_ci}')
    print(f'Male {confidence_level*100}% range:{male_range}')
    print(f'Female {confidence_level*100}% confidence interval:{female_ci}')
    print(f'Female {confidence_level*100}% range:{female_range}')
    print('\n')

```

## 5.1 Finding the Confidence Interval of the Age column for the whole dataset.

### 5.1.1 Analysing Gender across different and Confidence Interval to have a conclusive understanding the Spending pattern

#### *90% Confidence Interval*

```
df_ci('Gender', 0.90)
```

Records of Gender column

Male 90.0% confidence interval:[9424.512497305488, 9450.539583639042]

Male 90.0% range:26.027086333553598

Female 90.0% confidence interval:[8713.287834648021, 8755.84369566293]

Female 90.0% range:42.55586101490917

### ***95% Confidence Interval***

```
df_ci('Gender', 0.95)
```

Records of Gender column

Male 95.0% confidence interval:[9422.01944736257, 9453.032633581959]

Male 95.0% range:31.013186219388444

Female 95.0% confidence interval:[8709.21154714068, 8759.919983170272]

Female 95.0% range:50.70843602959212

### ***99% Confidence Interval***

```
df_ci('Gender', 0.99)
```

Records of Gender column

Male 99.0% confidence interval:[9417.146922669479, 9457.90515827505]

Male 99.0% range:40.75823560557183

Female 99.0% confidence interval:[8701.244674438389, 8767.886855872563]

Female 99.0% range:66.64218143417384

## **5.2 Finding the Confidence Interval of the Gender column across using different sample sizes.**

### **5.2.1 Analysing Gender column across different sample size and Confidence Interval to have a conclusive understanding the Spending pattern.**

#### **90% Confidence Interval**

*Sample Size: 300*

```
sample_ci('Gender', 0.90, 300)
```

300 samples of Gender column  
Male 90.0% confidence interval:[9380.83746339652, 10394.609203270147]  
Male 90.0% range:1013.7717398736277  
Female 90.0% confidence interval:[8398.670945742073, 9336.229054257929]  
Female 90.0% range:937.5581085158556

*Sample Size: 3000*

```
sample_ci('Gender', 0.90, 3000)
```

3000 samples of Gender column  
Male 90.0% confidence interval:[9489.930814851989, 9801.341185148012]  
Male 90.0% range:311.41037029602376  
Female 90.0% confidence interval:[8658.782721289897, 8954.244612043434]  
Female 90.0% range:295.46189075353686

*Sample Size: 30000*

```
sample_ci('Gender', 0.90, 30000)
```

30000 samples of Gender column  
Male 90.0% confidence interval:[9438.269538879862, 9535.561994453472]  
Male 90.0% range:97.29245557360991  
Female 90.0% confidence interval:[8610.964490517577, 8700.998442815755]  
Female 90.0% range:90.03395229817761

## 95% Confidence Interval

*Sample Size: 300*

```
sample_ci('Gender', 0.95, 300)
```

300 samples of Gender column  
Male 95.0% confidence interval:[9283.731565877591, 10491.715100789075]  
Male 95.0% range:1207.9835349114837  
Female 95.0% confidence interval:[8308.865304074718, 9426.034695925284]  
Female 95.0% range:1117.1693918505662

*Sample Size: 3000*

```
sample_ci('Gender', 0.95, 3000)
```



3000 samples of Gender column  
Male 95.0% confidence interval:[9460.10182838994, 9831.170171610062]  
Male 95.0% range:371.0683432201222  
Female 95.0% confidence interval:[8630.48138780842, 8982.545945524911]  
Female 95.0% range:352.0645577164905

*Sample Size: 30000*

```
sample_ci('Gender', 0.95, 30000)
```

30000 samples of Gender column  
Male 95.0% confidence interval:[9428.950211018666, 9544.881322314668]  
Male 95.0% range:115.9311112960022  
Female 95.0% confidence interval:[8602.340431075772, 8709.62250225756]  
Female 95.0% range:107.28207118178761

## 99% Confidence Interval

*Sample Size: 300*

```
sample_ci('Gender', 0.99, 300)
```

300 samples of Gender column  
Male 99.0% confidence interval:[9093.943597589929, 10681.503069076738]  
Male 99.0% range:1587.5594714868093  
Female 99.0% confidence interval:[8133.345271735851, 9601.55472826415]  
Female 99.0% range:1468.2094565282978

*Sample Size: 3000*

```
sample_ci('Gender', 0.99, 3000)
```

3000 samples of Gender column  
Male 99.0% confidence interval:[9401.80276642402, 9889.469233575981]  
Male 99.0% range:487.6664671519611  
Female 99.0% confidence interval:[8575.168036935633, 9037.859296397699]  
Female 99.0% range:462.6912594620662

*Sample Size: 30000*

```
sample_ci('Gender', 0.99, 30000)
```

30000 samples of Gender column

Male 99.0% confidence interval:[9410.736113360614, 9563.09541997272]

Male 99.0% range:152.3593066121066

Female 99.0% confidence interval:[8585.485196103737, 8726.477737229596]

Female 99.0% range:140.99254112585913

*Insights* - The analysis emphasizes the significance of sample size in determining population parameters. It indicates that larger sample sizes lead to narrower and more precise confidence intervals. In a business context, this means that bigger sample sizes can yield more dependable insights and estimates. - There are overlaps found in the confidence intervals.